

01_predictive_modelling_1

October 12, 2021

1 Wells Fargo Challenge

- <https://www.mindsumo.com/contests/campus-analytics-challenge-2021>

1.0.1 To Complete a Submission:

Build a classification model for predicting elder fraud in the digital payments space as described in Rule 4, which:

- Handles missing variables
- Maximizes the F1 score
- Uses the given data set
- Includes suitable encoding schemes
- Has the least set of feature variables

1.0.2 Resources

- <https://github.com/pdglenn/WellsFargoAnalyticsChallenge>

```
[1]: import pandas as pd
import numpy as np
import pylab as plt
import seaborn as sns

data_dir = "./dataset/"

# following few lines are to suppress the pandas warnings
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)

pd.options.mode.chained_assignment = None
pd.options.display.max_columns = 20
np.set_printoptions(suppress=True)

data_dir = "./dataset/"
image_dir = "./images/"
```

1.1 Loading the data

Note `pd.read_excel` gave me an error while reading the `xlsx` file so had to install `openpyxl` using `pip3 install openpyxl` and give `engine=openpyxl` as an extra argument.

```
[2]: #!pip3 install openpyxl
```

```
[3]: # load the file
df_orig = pd.read_excel(data_dir+"trainset.xlsx", engine='openpyxl')
df_orig.head(2)
```

```
[3]:   TRAN_AMT  ACCT_PRE_TRAN_AVAIL_BAL  CUST_AGE  OPEN_ACCT_CT  WF_dvc_age  \
0      5.38                23619.91        47             4      2777
1     65.19                 0.00        45             5      2721
```

```
      PWD_UPDT_TS      CARR_NAME  RGN_NAME STATE_PRVNC_TXT  \
0  1/16/2018 11:3:58  cox communications inc.  southwest      nevada
1           NaN  charter communications  southwest      california
```

```
      ALERT_TRGR_CD  ... CUST_STATE      PH_NUM_UPDT_TS  CUST_SINCE_DT  \
0           MOBL  ...      NV  2/24/2021 15:55:10      1993-01-06
1           MOBL  ...      CA           NaN      1971-01-07
```

```
      TRAN_TS      TRAN_DT  ACTN_CD  ACTN_INTNL_TXT  TRAN_TYPE_CD  \
0   5/3/2021 18:3:58   5/3/2021  SCHPMT      P2P_COMMIT      P2P
1  1/13/2021 19:19:37  1/13/2021  SCHPMT      P2P_COMMIT      P2P
```

```
      ACTVY_DT  FRAUD_NONFRAUD
0   5/3/2021      Non-Fraud
1  1/13/2021      Non-Fraud
```

[2 rows x 24 columns]

```
[4]: print ("Original data shape:", df_orig.shape)
```

Original data shape: (14000, 24)

```
[5]: #information of the dataset
df_orig.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14000 entries, 0 to 13999
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   TRAN_AMT                             14000 non-null  float64
1   ACCT_PRE_TRAN_AVAIL_BAL              14000 non-null  float64
2   CUST_AGE                             14000 non-null  int64
```

```

3  OPEN_ACCT_CT          14000 non-null  int64
4  WF_dvc_age           14000 non-null  int64
5  PWD_UPDT_TS          10875 non-null  object
6  CARR_NAME            11291 non-null  object
7  RGN_NAME             11291 non-null  object
8  STATE_PRVNC_TXT      11291 non-null  object
9  ALERT_TRGR_CD        14000 non-null  object
10 DVC_TYPE_TXT          12239 non-null  object
11 AUTHC_PRIM_TYPE_CD   14000 non-null  object
12 AUTHC_SCNDRY_STAT_TXT 13926 non-null  object
13 CUST_ZIP             14000 non-null  int64
14 CUST_STATE           13964 non-null  object
15 PH_NUM_UPDT_TS       6939 non-null  object
16 CUST_SINCE_DT        14000 non-null  datetime64[ns]
17 TRAN_TS             14000 non-null  object
18 TRAN_DT             14000 non-null  object
19 ACTN_CD             14000 non-null  object
20 ACTN_INTNL_TXT       14000 non-null  object
21 TRAN_TYPE_CD         14000 non-null  object
22 ACTVY_DT            14000 non-null  object
23 FRAUD_NONFRAUD       14000 non-null  object
dtypes: datetime64[ns](1), float64(2), int64(4), object(17)
memory usage: 2.6+ MB

```

```
[6]: # check the target classes
df_orig["FRAUD_NONFRAUD"].unique()
```

```
[6]: array(['Non-Fraud', 'Fraud'], dtype=object)
```

1.2 Train test split

Before doing any data visualization let's set some test data aside and use them to score the model later on.

```
[7]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# stratify the target column so that the distribution look similar in the train_
↪ and test data
df_train0, df_test0 = train_test_split(df_orig,
                                       test_size = .2,
                                       random_state = 8848,
                                       shuffle = True,
                                       stratify = df_orig["FRAUD_NONFRAUD"])
```

```
[8]: df = df_train0.copy()
```

```
[9]: df.head(2)
```

```
[9]:
```

	TRAN_AMT	ACCT_PRE_TRAN_AVAIL_BAL	CUST_AGE	OPEN_ACCT_CT	WF_dvc_age	\
2413	487.93	3714.91	43	5	1037	
1003	4.84	0.00	53	5	305	

	PWD_UPDT_TS	CARR_NAME	RGN_NAME	STATE_PRVNC_TXT	ALERT_TRGR_CD	\
2413	NaN	NaN	NaN	NaN	MOBL	
1003	4/12/2017 15:54:53	NaN	NaN	NaN	MOBL	

	...	CUST_STATE	PH_NUM_UPDT_TS	CUST_SINCE_DT	TRAN_TS	\
2413	...	CO	5/0/2020 12:33:41	1988-01-11	4/13/2021 5:2:29	
1003	...	TX	NaN	1987-04-05	4/29/2021 22:54:53	

	TRAN_DT	ACTN_CD	ACTN_INTNL_TXT	TRAN_TYPE_CD	ACTVY_DT	FRAUD_NONFRAUD
2413	4/13/2021	SCHPMT	P2P_COMMIT	P2P	4/13/2021	Fraud
1003	4/29/2021	SCHPMT	P2P_COMMIT	P2P	4/29/2021	Non-Fraud

[2 rows x 24 columns]

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 11200 entries, 2413 to 114
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   TRAN_AMT                             11200 non-null  float64
1   ACCT_PRE_TRAN_AVAIL_BAL              11200 non-null  float64
2   CUST_AGE                             11200 non-null  int64
3   OPEN_ACCT_CT                         11200 non-null  int64
4   WF_dvc_age                           11200 non-null  int64
5   PWD_UPDT_TS                          8684 non-null   object
6   CARR_NAME                            9022 non-null   object
7   RGN_NAME                             9022 non-null   object
8   STATE_PRVNC_TXT                      9022 non-null   object
9   ALERT_TRGR_CD                        11200 non-null  object
10  DVC_TYPE_TXT                          9805 non-null   object
11  AUTHC_PRIM_TYPE_CD                   11200 non-null  object
12  AUTHC_SCNDRY_STAT_TXT                11140 non-null  object
13  CUST_ZIP                             11200 non-null  int64
14  CUST_STATE                           11172 non-null  object
15  PH_NUM_UPDT_TS                       5579 non-null   object
16  CUST_SINCE_DT                        11200 non-null  datetime64[ns]
17  TRAN_TS                              11200 non-null  object
18  TRAN_DT                              11200 non-null  object
19  ACTN_CD                              11200 non-null  object
20  ACTN_INTNL_TXT                       11200 non-null  object
21  TRAN_TYPE_CD                         11200 non-null  object
```

```

22  ACTVY_DT                11200 non-null object
23  FRAUD_NONFRAUD          11200 non-null object
dtypes: datetime64[ns](1), float64(2), int64(4), object(17)
memory usage: 2.1+ MB

```

```
[11]: df.shape
```

```
[11]: (11200, 24)
```

```
[12]: # find numerical and categorical columns
nume_cols = list(df.select_dtypes(include="number").columns)
cate_cols = list(df.select_dtypes(exclude="number").columns)
cate_cols.remove('FRAUD_NONFRAUD')
```

```
[13]: print ("Numeric Columns:\n", nume_cols)
print ("")
print ("Categorical Columns:\n", cate_cols)
```

Numeric Columns:

```
['TRAN_AMT', 'ACCT_PRE_TRAN_AVAIL_BAL', 'CUST_AGE', 'OPEN_ACCT_CT',
'WF_dvc_age', 'CUST_ZIP']
```

Categorical Columns:

```
['PWD_UPDT_TS', 'CARR_NAME', 'RGN_NAME', 'STATE_PRVNC_TXT', 'ALERT_TRGR_CD',
'DVC_TYPE_TXT', 'AUTHC_PRIM_TYPE_CD', 'AUTHC_SCNDRY_STAT_TXT', 'CUST_STATE',
'PH_NUM_UPDT_TS', 'CUST_SINCE_DT', 'TRAN_TS', 'TRAN_DT', 'ACTN_CD',
'ACTN_INTNL_TXT', 'TRAN_TYPE_CD', 'ACTVY_DT']
```

```
[14]: df[nume_cols].head(2)
```

```
[14]:
```

	TRAN_AMT	ACCT_PRE_TRAN_AVAIL_BAL	CUST_AGE	OPEN_ACCT_CT	WF_dvc_age	\
2413	487.93	3714.91	43	5	1037	
1003	4.84	0.00	53	5	305	

	CUST_ZIP
2413	80234
1003	75232

```
[15]: nume_cols.remove('CUST_ZIP')
cate_cols.append('CUST_ZIP')
```

```
[16]: print ("Numeric Columns:\n", nume_cols)
print ("")
print ("Categorical Columns:\n", cate_cols)
```

Numeric Columns:

```
['TRAN_AMT', 'ACCT_PRE_TRAN_AVAIL_BAL', 'CUST_AGE', 'OPEN_ACCT_CT',
'WF_dvc_age']
```

Categorical Columns:

```
['PWD_UPDT_TS', 'CARR_NAME', 'RGN_NAME', 'STATE_PRVNC_TXT', 'ALERT_TRGR_CD',  
'DVC_TYPE_TXT', 'AUTHC_PRIM_TYPE_CD', 'AUTHC_SCNDRY_STAT_TXT', 'CUST_STATE',  
'PH_NUM_UPDT_TS', 'CUST_SINCE_DT', 'TRAN_TS', 'TRAN_DT', 'ACTN_CD',  
'ACTN_INTNL_TXT', 'TRAN_TYPE_CD', 'ACTVY_DT', 'CUST_ZIP']
```

```
[17]: df[nume_cols].head(5)
```

```
[17]:
```

	TRAN_AMT	ACCT_PRE_TRAN_AVAIL_BAL	CUST_AGE	OPEN_ACCT_CT	WF_dvc_age
2413	487.93	3714.91	43	5	1037
1003	4.84	0.00	53	5	305
8660	494.94	2525.50	70	9	583
6349	0.01	0.00	70	6	467
1860	488.36	4344.55	38	4	0

```
[18]: impute_vals={}

for col in df.columns:
    if col in nume_cols:
        impute_vals[col] = df[col].median()
    elif col in cate_cols:
        impute_vals[col] = df[col].mode()[0]

impute_vals
```

```
[18]: {'TRAN_AMT': 162.07,  
      'ACCT_PRE_TRAN_AVAIL_BAL': 2396.1549999999997,  
      'CUST_AGE': 59.0,  
      'OPEN_ACCT_CT': 5.0,  
      'WF_dvc_age': 366.5,  
      'PWD_UPDT_TS': '5/18/2020 4:7:20',  
      'CARR_NAME': 'cox communications inc.',  
      'RGN_NAME': 'southwest',  
      'STATE_PRVNC_TXT': 'california',  
      'ALERT_TRGR_CD': 'MOBL',  
      'DVC_TYPE_TXT': 'MOBILE',  
      'AUTHC_PRIM_TYPE_CD': 'UN_PWD',  
      'AUTHC_SCNDRY_STAT_TXT': 'ALLOW',  
      'CUST_ZIP': 77459,  
      'CUST_STATE': 'CA',  
      'PH_NUM_UPDT_TS': '7/8/2019 6:45:37',  
      'CUST_SINCE_DT': Timestamp('1997-08-01 00:00:00'),  
      'TRAN_TS': '2/3/2021 9:0:51',  
      'TRAN_DT': '2/28/2021',  
      'ACTN_CD': 'SCHPMT',  
      'ACTN_INTNL_TXT': 'P2P_COMMIT',
```

```
'TRAN_TYPE_CD': 'P2P',
'ACTVY_DT': '2/28/2021'}
```

```
[19]: cols_to_drop = ['PH_NUM_UPDT_TS', 'PWD_UPDT_TS', 'CARR_NAME', 'RGN_NAME',
    ↪ 'STATE_PRVNC_TXT']
nume_cols = [c for c in nume_cols if c not in cols_to_drop]
cate_cols = [c for c in cate_cols if c not in cols_to_drop]
```

```
[20]: def impute_data(df, impute_dict=impute_vals):
    """
    this function takes in a dataframe and list of columns which have missing_
    ↪ values
    then imputes those columns using the precomputed values.
    """
    for col in list(impute_dict.keys()):
        df[col] = df[col].fillna(impute_dict[col])
    return df
```

```
[21]: # impute the columns : cols_to_impute
df=impute_data(df)
```

```
[22]: df.isnull().sum()
```

```
[22]: TRAN_AMT                                0
ACCT_PRE_TRAN_AVAIL_BAL                      0
CUST_AGE                                     0
OPEN_ACCT_CT                                0
WF_dvc_age                                  0
PWD_UPDT_TS                                 0
CARR_NAME                                    0
RGN_NAME                                    0
STATE_PRVNC_TXT                             0
ALERT_TRGR_CD                               0
DVC_TYPE_TXT                                0
AUTHC_PRIM_TYPE_CD                          0
AUTHC_SCNDRY_STAT_TXT                       0
CUST_ZIP                                     0
CUST_STATE                                  0
PH_NUM_UPDT_TS                              0
CUST_SINCE_DT                               0
TRAN_TS                                     0
TRAN_DT                                     0
ACTN_CD                                     0
ACTN_INTNL_TXT                              0
TRAN_TYPE_CD                                0
ACTVY_DT                                    0
FRAUD_NONFRAUD                              0
```

dtype: int64

```
[23]: df[nume_cols].head(2)
```

```
[23]:
```

	TRAN_AMT	ACCT_PRE_TRAN_AVAIL_BAL	CUST_AGE	OPEN_ACCT_CT	WF_dvc_age
2413	487.93	3714.91	43	5	1037
1003	4.84	0.00	53	5	305

1.3 Feature Transformation

- Convert the categorical features into small number of categories when possible

```
[24]: def wrangle_column_data(df):  
    #CUST_STATE  
    # keep only CO, TX, MN, AZ and convert rest into OTHER  
    df["CUST_STATE"] = df["CUST_STATE"].apply(lambda x: x if x in ["CO", "TX",  
    ↪ "MN", "AZ"] else "OTHER")  
    #OPEN_ACCT_CT  
    #keep only the [2,9] and convert others to 10  
    df["OPEN_ACCT_CT"] = df["OPEN_ACCT_CT"].apply(lambda x: x if x in  
    ↪ range(1,13) else 13)  
    return df
```

```
[25]: df=wrangle_column_data(df)
```

```
[26]: nume_cols
```

```
[26]: ['TRAN_AMT',  
      'ACCT_PRE_TRAN_AVAIL_BAL',  
      'CUST_AGE',  
      'OPEN_ACCT_CT',  
      'WF_dvc_age']
```

```
[27]: cate_cols_to_keep = ['ALERT_TRGR_CD', "CUST_STATE"]
```

```
[28]: df[cate_cols_to_keep].head()
```

```
[28]:
```

	ALERT_TRGR_CD	CUST_STATE
2413	MOBL	CO
1003	MOBL	TX
8660	MOBL	TX
6349	ONLN	MN
1860	MOBL	AZ

1.4 Build a model with only Numerical features

```
[29]: # map Fraud to 1 and Non-Fraud to 0 in the dataframe for both train and test_
      ↪data
df["FRAUD_NONFRAUD"] = df["FRAUD_NONFRAUD"].map({"Fraud":1, "Non-Fraud":0})
```

```
[30]: X_train1, y_train1 = df[nume_cols], df["FRAUD_NONFRAUD"]
```

```
[31]: # prepare test data
df_test = df_test0.copy()
print ("missing values in test data:\n", df_test[nume_cols].isnull().sum() )
df_test["FRAUD_NONFRAUD"] = df_test["FRAUD_NONFRAUD"].map({"Fraud":
      ↪1, "Non-Fraud":0})
X_test1, y_test1 = df_test[nume_cols], df_test["FRAUD_NONFRAUD"]
```

```
missing values in test data:
  TRAN_AMT          0
ACCT_PRE_TRAN_AVAIL_BAL  0
  CUST_AGE          0
  OPEN_ACCT_CT        0
  WF_dvc_age          0
dtype: int64
```

```
[32]: X_train1.shape, y_train1.shape, X_test1.shape, y_test1.shape
```

```
[32]: ((11200, 5), (11200,), (2800, 5), (2800,))
```

1.4.1 Base Model: Logistic Regression, Random Forest, XGBoost

```
[33]: from sklearn.linear_model import LogisticRegression
      from sklearn.model_selection import GridSearchCV
      from sklearn.ensemble import RandomForestClassifier, VotingClassifier
      from xgboost import XGBClassifier
      from sklearn.metrics import accuracy_score, f1_score, precision_score,
      ↪recall_score
      from sklearn.metrics import classification_report
      from sklearn.metrics import roc_auc_score, plot_roc_curve
```

```
[34]: class Model_training:
      def __init__(self, model, X_train, y_train, X_test, y_test, savename="Fig"):
          self.model = model
          self.X_train = X_train
          self.y_train = y_train
          self.X_test = X_test
          self.y_test = y_test
          self.savename = savename
```

```

self.model.fit(self.X_train, self.y_train)

def print_metrics(self):
    round_to_pct = lambda x: np.round(100*x, 2)
    y_pred = self.model.predict(self.X_test)
    ac = round_to_pct(accuracy_score(self.y_test, y_pred))
    f1 = round_to_pct(f1_score(self.y_test, y_pred))
    pr = round_to_pct(precision_score(self.y_test, y_pred))
    re = round_to_pct(recall_score(self.y_test, y_pred))
    print (f"Accuracy = {ac}% F1 Score= {f1}% \nPrecision={pr}% Recall=
→{re}%")
    print (classification_report(self.y_test, y_pred))
    return (self.model, (ac, f1, pr, re))

def displot(self):
    pr=self.model.predict_proba(self.X_test)
    roc_auc = np.round(roc_auc_score(self.y_test,
                                     self.model.predict_proba(self.X_test)[:
→, 1]), 2)

    pr_df = pd.DataFrame({'pred_0':pr[:,0],
                          'pred_1':pr[:,1],
                          'y': self.y_test})

    ax=sns.displot(data=pr_df,
                   x='pred_1',
                   hue='y',
                   alpha=0.8,
                   kind="kde",
                   height = 3.5,
                   aspect=1.8);

    plt.xlabel("Prob. Positive Predictions", fontsize=16)
    plt.text(0.2, 2, "ROC_AUC="+str(roc_auc), fontsize=16)
    plt.ylabel("Density", fontsize=16)

    plt.yticks(fontsize=16);
    plt.xticks(fontsize=16);
    figname = "images/displot_"+self.savename+"_nb1.png"
    plt.savefig(figname, dpi=300, bbox_inches='tight')

def feature_importance(self):

    df_imp = pd.DataFrame({"Feature":self.X_train.columns,
                          "Feature Importance":self.model.
→feature_importances_})

```

```

df_imp = df_imp.sort_values(by=['Feature Importance'],
                             axis=0,
                             ascending=True)

df_imp.plot(kind='barh',
            x='Feature',
            y='Feature Importance',
            color="C2", figsize=(8,5));

plt.grid(axis='x')
plt.yticks(fontsize=16);
plt.ylabel('');
plt.xticks(fontsize=16);
plt.legend(loc='best',fontsize=16);

figname = "images/feat_imp_"+self.savename+"_nb1.png"
plt.savefig(figname, dpi=300, bbox_inches='tight')

def plot_roc_curve(self):
    roc_auc = np.round(roc_auc_score(self.y_test,
                                     self.model.predict_proba(self.X_test)[:
→, 1]), 2)
    label_name = self.savename + "\n AUC = "+str(roc_auc)

    plot_roc_curve(self.model, self.X_test, self.y_test,
                  lw=3., color='C2', label=label_name)
    plt.xlabel("False Positive Rate", fontsize=16)
    plt.ylabel("True Positive Rate", fontsize=16)
    plt.xticks(fontsize=16);
    plt.yticks(fontsize=16);
    plt.legend(loc="center", fontsize=14);
    plt.axvline(x=0, color='k', ls='--', lw=1)
    plt.axhline(y=0, color='k', ls='--', lw=1)
    plt.axhline(y=1, color='k', ls='--', lw=1)

    figname = "images/roc_curve_"+self.savename+"_nb1.png"
    plt.savefig(figname, dpi=300, bbox_inches='tight')

```

```

[35]: model_rf = RandomForestClassifier(max_depth=10,
                                     random_state=8848)

mod = Model_training(model_rf,
                    X_train1, y_train1,
                    X_test1, y_test1,
                    "Random_Forest")
mod_tr, _ = mod.print_metrics()
mod.displot()


```



```
mod.plot_roc_curve()  
mod.feature_importance()
```

Accuracy = 93.64% F1 Score= 88.81%

Precision=93.26% Recall= 84.75%

	precision	recall	f1-score	support
0	0.94	0.97	0.96	1967
1	0.93	0.85	0.89	833
accuracy			0.94	2800
macro avg	0.94	0.91	0.92	2800
weighted avg	0.94	0.94	0.94	2800





For a base model: F1 score ~ 88% is a great result.

```
[36]: model_rf_gs = GridSearchCV(RandomForestClassifier(),
                                param_grid={'max_depth':[8, 9, 10, 11, 12, 14]},
                                scoring = 'f1',
                                verbose=1)
```

```
mod2 = Model_training(model_rf_gs,
                      X_train1, y_train1, X_test1, y_test1,
                      "random_forest_grid_search")
mod_tr, _ = mod2.print_metrics()
mod2.displot()
```

Fitting 5 folds for each of 6 candidates, totalling 30 fits


[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 32.0s finished

Accuracy = 93.57% F1 Score= 88.78%

Precision=92.35% Recall= 85.47%

	precision	recall	f1-score	support
0	0.94	0.97	0.95	1967
1	0.92	0.85	0.89	833
accuracy			0.94	2800
macro avg	0.93	0.91	0.92	2800
weighted avg	0.94	0.94	0.93	2800



```
[37]: mod_tr.best_params_
```

```
[37]: {'max_depth': 14}
```

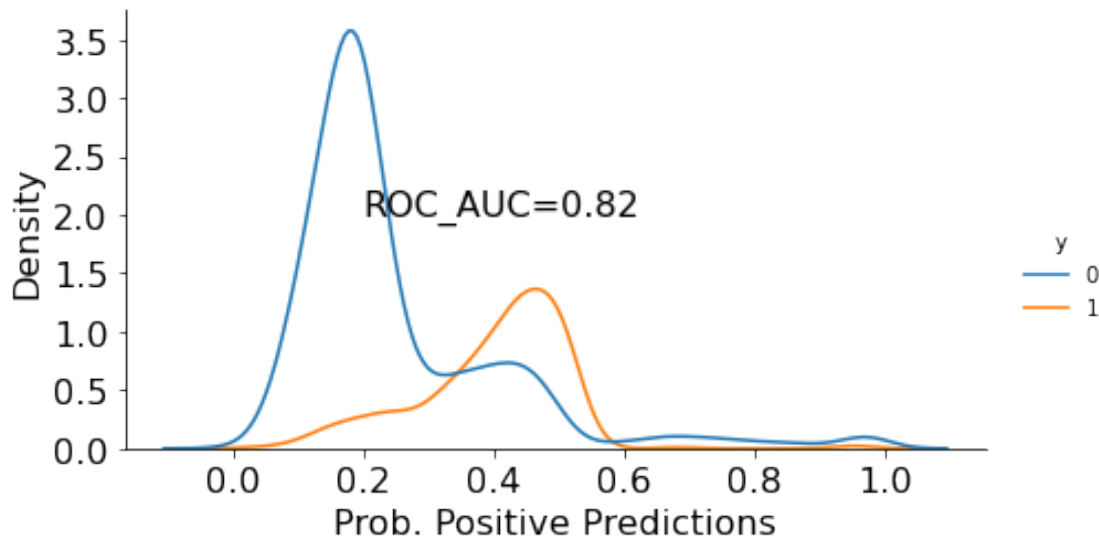
```
[38]: model_lr = LogisticRegression(max_iter=5000)
```

```
mod3 = Model_training(model_lr, X_train1, y_train1, X_test1, y_test1,
↳ "logistic_regression")
mod_tr, _ = mod3.print_metrics()
mod3.displot()
```

Accuracy = 70.86% F1 Score= 23.02%

Precision=53.74% Recall= 14.65%

	precision	recall	f1-score	support
0	0.72	0.95	0.82	1967
1	0.54	0.15	0.23	833
accuracy			0.71	2800
macro avg	0.63	0.55	0.53	2800
weighted avg	0.67	0.71	0.64	2800





```
[39]: xgb = XGBClassifier(verbosity=1,
                        use_label_encoder=False,
                        max_depth=10,
                        eval_metric = "logloss")

mod4 = Model_training(xgb, X_train1, y_train1, X_test1, y_test1, "xgb")
mod_tr, _ = mod4.print_metrics()
mod4.displot()
mod4.feature_importance()
```

Accuracy = 93.14% F1 Score= 88.13%

Precision=90.83% Recall= 85.59%

	precision	recall	f1-score	support
0	0.94	0.96	0.95	1967
1	0.91	0.86	0.88	833
accuracy			0.93	2800
macro avg	0.92	0.91	0.92	2800
weighted avg	0.93	0.93	0.93	2800





```
[40]: xgb_gs = GridSearchCV(XGBClassifier(),
                             param_grid={'max_depth':[8, 9, 10, 11, 12, 14],
                                           'eval_metric':["logloss"],
                                           'reg_alpha':[0.1, 0.5]},
                             scoring = 'f1',
                             verbose = 0 )

mod5 = Model_training(xgb_gs, X_train1, y_train1, X_test1, y_test1,
↳"xgb_grid_search")
mod_tr, _ = mod5.print_metrics()
mod5.displot()
```

Accuracy = 92.93% F1 Score= 87.78%

Precision=90.34% Recall= 85.35%

	precision	recall	f1-score	support
0	0.94	0.96	0.95	1967
1	0.90	0.85	0.88	833
accuracy			0.93	2800
macro avg	0.92	0.91	0.91	2800
weighted avg	0.93	0.93	0.93	2800



```
[41]: mod_tr.best_params_
```

```
[41]: {'eval_metric': 'logloss', 'max_depth': 11, 'reg_alpha': 0.1}
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

1.5 Deep learning models

```
[42]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.optimizers import Adam
      from tensorflow.keras.layers import Dense, Dropout, Flatten
      from tensorflow.keras.layers import Conv1D, MaxPooling1D, GlobalMaxPool1D,
      ↪BatchNormalization
      from tensorflow.keras import backend as K

[43]: # this piece of code copied from
      #https://neptune.ai/blog/implementing-the-macro-f1-score-in-keras
      def custom_f1(y_true, y_pred):
          def recall_m(y_true, y_pred):
              TP = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
              Positives = K.sum(K.round(K.clip(y_true, 0, 1)))

              recall = TP / (Positives+K.epsilon())
              return recall

          def precision_m(y_true, y_pred):
              TP = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
              Pred_Positives = K.sum(K.round(K.clip(y_pred, 0, 1)))

              precision = TP / (Pred_Positives+K.epsilon())
              return precision

          precision, recall = precision_m(y_true, y_pred), recall_m(y_true, y_pred)

          return 2*((precision*recall)/(precision+recall+K.epsilon()))

[44]: def make_model_dense(X_train, y_train):
      model = Sequential()
      model.add(Dense(1, input_shape=(X_train.shape[1],), activation='relu'))
      model.add(Dropout(0.25))
      model.add(Dense(1, activation='relu'))
      opt = Adam(learning_rate=0.001)
      model.compile(loss='binary_crossentropy',
                    optimizer=opt,
```

```
        metrics=custom_f1)
    return model
```

```
[45]: model_dnn = make_model_dense(X_train1, y_train1)
      model_dnn.summary()
```

Model: "sequential"


Layer (type)	Output Shape	Param #
dense (Dense)	(None, 1)	6
dropout (Dropout)	(None, 1)	0
dense_1 (Dense)	(None, 1)	2

Total params: 8
Trainable params: 8
Non-trainable params: 0

```
[46]: history_dnn = model_dnn.fit(X_train1, y_train1,
                                validation_data=(X_test1, y_test1),
                                epochs=20,
                                batch_size=32,
                                verbose=0)
```

```
[47]: plt.plot(history_dnn.history['val_custom_f1'])
      plt.plot(history_dnn.history['val_loss'])
```

```
[47]: [<matplotlib.lines.Line2D at 0x158570490>]
```



1.6 Modeling including categorical features

```
[48]: import category_encoders as ce
```

```
[49]: df2 = df_train0.copy()
df2["FRAUD_NONFRAUD"] = df2["FRAUD_NONFRAUD"].map({"Fraud":1, "Non-Fraud":0})
df2 = impute_data(df2)
df2 = wrangle_column_data(df2)
```

```
[50]: encoder = ce.OneHotEncoder()
df2_tr_cat = encoder.fit_transform(df[cate_cols_to_keep])
df2_tr_join = pd.concat([df2[nume_cols], df2_tr_cat], axis=1)
X_train2 = df2_tr_join.values
y_train2 = df2["FRAUD_NONFRAUD"].values
```

```
[51]: # test data

df2_te = df_test0.copy()
df2_te["FRAUD_NONFRAUD"] = df2_te["FRAUD_NONFRAUD"].map({"Fraud":1, "Non-Fraud":
    ↪0})
df2_te = impute_data(df2_te)
df2_te = wrangle_column_data(df2_te)
```

```
[52]: df2_te_cat = encoder.transform(df2_te[cate_cols_to_keep])
df2_te_join = pd.concat([df2_te[nume_cols], df2_te_cat], axis=1)
```

```
X_test2 = df2_te_join.values
y_test2 = df2_te["FRAUD_NONFRAUD"].values
```

```
[53]: X_train2.shape, y_train2.shape, X_test2.shape, y_test2.shape
```

```
[53]: ((11200, 12), (11200,), (2800, 12), (2800,))
```


```
[54]: xgb = XGBClassifier(verbosity=1,
                        max_depth=10,
                        eval_metric = "logloss")

mod6 = Model_training(xgb, X_train2, y_train2, X_test2, y_test2)
mod_trained, _ = mod5.print_metrics()
mod6.displot()
```

Accuracy = 92.93% F1 Score= 87.78%

Precision=90.34% Recall= 85.35%

	precision	recall	f1-score	support
0	0.94	0.96	0.95	1967
1	0.90	0.85	0.88	833
accuracy			0.93	2800
macro avg	0.92	0.91	0.91	2800
weighted avg	0.93	0.93	0.93	2800



```
[ ]:
```


```
[55]: from sklearn.svm import SVC
```

```
[56]: svm = SVC(gamma='auto')
mod = Model_training(xgb_gs, X_train1, y_train1, X_test1, y_test1)
mod_tr, _ = mod.print_metrics()
mod.displot()
```

Accuracy = 92.93% F1 Score= 87.78%

Precision=90.34% Recall= 85.35%

	precision	recall	f1-score	support
0	0.94	0.96	0.95	1967
1	0.90	0.85	0.88	833
accuracy			0.93	2800
macro avg	0.92	0.91	0.91	2800
weighted avg	0.93	0.93	0.93	2800



```
[57]: def train_svm_gs():
    svm_gs = GridSearchCV(SVC(),
                           param_grid={'C':[0.1, 0.5, 1.0],
                                         'kernel':['poly', 'rbf'],
                                         'gamma': ['scale', 'auto']},
                           scoring = 'f1',
                           verbose = 1 )

    mod = Model_training(svm_gs, X_train1, y_train1, X_test1, y_test1)
```

```
mod_tr, _ = mod.print_metrics()
mod.displot()
```

#this takes a little long time so think before running
#train_sum_gs()

1.7 Voting Classifier

```
[58]: clf1 = LogisticRegression(multi_class='multinomial', random_state=1)
      clf2 = RandomForestClassifier(max_depth=10,
                                   random_state=8848)
```

```
      clf3 = XGBClassifier(verbosity=1,
                           max_depth=13,
                           eval_metric = "logloss")
```



```
      clf_voting = VotingClassifier(
          estimators=[('rf', clf2), ('xgb', clf3)],
          voting='soft')
```

```
[59]: mod8 = Model_training(clf_voting, X_train1, y_train1, X_test1, y_test1, "Voting_
      ↪Classifier")
      mod_tr, _ = mod8.print_metrics()
      mod8.displot()
      mod8.plot_roc_curve()
      #mod8.feature_importance()
```

Accuracy = 93.46% F1 Score= 88.58%

Precision=92.21% Recall= 85.23%

	precision	recall	f1-score	support
0	0.94	0.97	0.95	1967
1	0.92	0.85	0.89	833
accuracy			0.93	2800
macro avg	0.93	0.91	0.92	2800
weighted avg	0.93	0.93	0.93	2800



[]:

[]: